

```
In [2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
In [3]: train = pd.read_csv('C:/Users/user/Downloads/titanic_train.csv')
```

```
In [4]: train.head()
```

```
Out[4]:
```

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...)	female	38.0	1	0	PC 17599	71.2
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0

```
In [5]: train.isnull()
```

Out[5]:


	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Ca
0	False	False	False	False	False	False	False	False	False	False	T
1	False	False	False	False	False	False	False	False	False	False	Fa
2	False	False	False	False	False	False	False	False	False	False	T
3	False	False	False	False	False	False	False	False	False	False	Fa
4	False	False	False	False	False	False	False	False	False	False	T
...	
886	False	False	False	False	False	False	False	False	False	False	T
887	False	False	False	False	False	False	False	False	False	False	Fa
888	False	False	False	False	False	True	False	False	False	False	T
889	False	False	False	False	False	False	False	False	False	False	Fa
890	False	False	False	False	False	False	False	False	False	False	T

891 rows × 12 columns




```
In [6]: sns.heatmap(train.isnull(),yticklabels=False,cbar=False,cmap='viridis')
```

Out[6]: <Axes: >




```
In [7]: sns.set_style('whitegrid')  
sns.countplot(x='Survived',data=train)
```

Out[7]: <Axes: xlabel='Survived', ylabel='count'>



```
In [8]: sns.set_style('whitegrid')  
sns.countplot(x='Survived',hue='Sex',data=train,palette='RdBu_r')
```


Out[8]: <Axes: xlabel='Survived', ylabel='count'>



```
In [9]: import seaborn as sns
import matplotlib.pyplot as plt

# Ensure the 'Pclass' column is of type string
train['Pclass'] = train['Pclass'].astype(str)

sns.set_style('whitegrid')
sns.countplot(x='Survived', hue='Pclass', data=train, palette='rainbow')
plt.show()
```



```
In [19]: sns.distplot(train['Age'].dropna(),kde=False,color='darkred',bins=40)
```

C:\Users\user\AppData\Local\Temp\ipykernel_23296\2002818437.py:1: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>


```
sns.distplot(train['Age'].dropna(),kde=False,color='darkred',bins=40)
```

```
Out[19]: <Axes: xlabel='Age'>
```




```
In [20]: sns.countplot(x='SibSp',data=train)
```

```
Out[20]: <Axes: xlabel='SibSp', ylabel='count'>
```



```
In [21]: train['Fare'].hist(color='green',bins=40,figsize=(8,4))
```


```
Out[21]: <Axes: >
```



```
In [ ]: #DATA CLEANING
```

```
In [4]: plt.figure(figsize=(12, 7))  
sns.boxplot(x='Pclass',y='Age',data=train,palette='winter')
```

```
Out[4]: <Axes: xlabel='Pclass', ylabel='Age'>
```



```
In [5]: def impute_age(cols):  
    Age = cols[0]  
    Pclass = cols[1]  
  
    if pd.isnull(Age):  
  
        if Pclass == 1:  
            return 37  
  
        elif Pclass == 2:  
            return 29
```

```

    else:
        return 24

    else:
        return Age

```

In []: *#Now apply that function!*

In [6]: `train['Age'] = train[['Age', 'Pclass']].apply(impute_age,axis=1)`

C:\Users\user\AppData\Local\Temp\ipykernel_16528\822839471.py:2: FutureWarning: Series.__getitem__ treating keys as positions is deprecated. In a future version, integer keys will always be treated as labels (consistent with DataFrame behavior). To access a value by position, use `ser.iloc[pos]`

Age = cols[0]


C:\Users\user\AppData\Local\Temp\ipykernel_16528\822839471.py:3: FutureWarning: Series.__getitem__ treating keys as positions is deprecated. In a future version, integer keys will always be treated as labels (consistent with DataFrame behavior). To access a value by position, use `ser.iloc[pos]`

Pclass = cols[1]

In []: *#Now Let's check that heat map again!*

In [7]: `sns.heatmap(train.isnull(),yticklabels=False,cbar=False,cmap='viridis')`

Out[7]: <Axes: >




In [60]: `train.drop('Cabin',axis=1,inplace=True)`
`train.head()`

Out[60]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0

In [61]: `train.dropna(inplace=True)`In [10]: `sns.heatmap(train.isnull(),yticklabels=False,cbar=False,cmap='viridis')`Out[10]: `<Axes: >`



```
In [11]: train.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Index: 889 entries, 0 to 890  
Data columns (total 11 columns):  
#   Column      Non-Null Count  Dtype  
---  ---  
0   PassengerId  889 non-null    int64  
1   Survived     889 non-null    int64  
2   Pclass       889 non-null    int64  
3   Name         889 non-null    object  
4   Sex          889 non-null    object  
5   Age          889 non-null    float64  
6   SibSp        889 non-null    int64  
7   Parch        889 non-null    int64  
8   Ticket       889 non-null    object  
9   Fare         889 non-null    float64  
10  Embarked     889 non-null    object  
dtypes: float64(2), int64(5), object(4)  
memory usage: 83.3+ KB
```

```
In [ ]: #We'll need to convert categorical features to dummy variables using pandas! Oth
```

```
In [12]: pd.get_dummies(train['Embarked'],drop_first=True).head()
```

```
Out[12]:
```

	Q	S
0	False	True
1	False	False
2	False	True
3	False	True
4	False	True

```
In [13]: sex = pd.get_dummies(train['Sex'],drop_first=True)
embark = pd.get_dummies(train['Embarked'],drop_first=True)
```

```
In [14]: train.drop(['Sex','Embarked','Name','Ticket'],axis=1,inplace=True)
```

```
In [15]: train.head()
```

```
Out[15]:
```

	PassengerId	Survived	Pclass	Age	SibSp	Parch	Fare
0	1	0	3	22.0	1	0	7.2500
1	2	1	1	38.0	1	0	71.2833
2	3	1	3	26.0	0	0	7.9250
3	4	1	1	35.0	1	0	53.1000
4	5	0	3	35.0	0	0	8.0500

```
In [16]: train = pd.concat([train,sex,embark],axis=1)
train.head()
```

```
Out[16]:
```

	PassengerId	Survived	Pclass	Age	SibSp	Parch	Fare	male	Q	S
0	1	0	3	22.0	1	0	7.2500	True	False	True
1	2	1	1	38.0	1	0	71.2833	False	False	False
2	3	1	3	26.0	0	0	7.9250	False	False	True
3	4	1	1	35.0	1	0	53.1000	False	False	True
4	5	0	3	35.0	0	0	8.0500	True	False	True

```
In [ ]: #LOGISTIC REGRESSION
```

```
In [ ]: #remove passenger id and name since it is unnecessary
```

```
In [17]: train.drop(['PassengerId'], axis=1, inplace=True)
train.head()
```

Out[17]:

	Survived	Pclass	Age	SibSp	Parch	Fare	male	Q	S
0	0	3	22.0	1	0	7.2500	True	False	True
1	1	1	38.0	1	0	71.2833	False	False	False
2	1	3	26.0	0	0	7.9250	False	False	True
3	1	1	35.0	1	0	53.1000	False	False	True
4	0	3	35.0	0	0	8.0500	True	False	True

In [18]: *# Convert boolean to integers (True -> 1, False -> 0)*
 train[['male', 'Q', 'S']] = train[['male', 'Q', 'S']].astype(int)

In [40]: *#Define Features and Target Variable*
#TRAIN TEST SPLIT

In [19]: train.drop('Survived',axis=1).head()

Out[19]:

	Pclass	Age	SibSp	Parch	Fare	male	Q	S
0	3	22.0	1	0	7.2500	1	0	1
1	1	38.0	1	0	71.2833	0	0	0
2	3	26.0	0	0	7.9250	0	0	1
3	1	35.0	1	0	53.1000	0	0	1
4	3	35.0	0	0	8.0500	1	0	1

In [20]: train['Survived'].head()

Out[20]: 0 0
 1 1
 2 1
 3 1
 4 0
 Name: Survived, dtype: int64

In [21]: *from* sklearn.model_selection *import* train_test_split
 X_train, X_test, y_train, y_test = train_test_split(train.drop('Survived',axis=1),
 train['Survived'], test_size
 random_state=101)

In []: *#TRAINING AND PREDICTING*

In [45]: *from* sklearn.linear_model *import* LogisticRegression

In [46]: logmodel = LogisticRegression()
 logmodel.fit(X_train,y_train)

D:\anaconda\Lib\site-packages\sklearn\linear_model_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

n_iter_i = _check_optimize_result(

Out[46]: ▾ LogisticRegression
LogisticRegression()

In [47]: predictions = logmodel.predict(X_test)

In []: #EVALUATION

In [48]: from sklearn.metrics import confusion_matrix

In [49]: accuracy=confusion_matrix(y_test,predictions)
accuracy

Out[49]: array([[147, 16],
[30, 74]], dtype=int64)

In [50]: from sklearn.metrics import accuracy_score
accuracy=accuracy_score(y_test,predictions)
accuracy

Out[50]: 0.8277153558052435

In [51]: predictions

Out[51]: array([0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 1, 1, 0, 1, 0, 0, 1, 1, 1, 0, 0, 0,
0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1,
1, 0, 1, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 0, 0, 0,
0, 0, 0, 0, 0, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1,
0, 1, 1, 1, 0, 0, 0, 1, 1, 0, 0, 1, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0,
0, 0, 1, 1, 0, 1, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 1, 0, 1, 1, 0, 0,
1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 1, 1,
0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0,
0, 0, 1, 0, 1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0,
0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0,
1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 1, 1, 1, 0, 1, 0, 0, 0,
0, 0, 1, 0, 0, 0, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0,
0, 1, 1], dtype=int64)

In [52]: from sklearn.metrics import classification_report
print(classification_report(y_test,predictions))

	precision	recall	f1-score	support
0	0.83	0.90	0.86	163
1	0.82	0.71	0.76	104
accuracy			0.83	267
macro avg	0.83	0.81	0.81	267
weighted avg	0.83	0.83	0.83	267

```
In [ ]: #RANDOM FOREST
```

```
In [35]: df = pd.read_csv('C:/Users/user/Downloads/titanic_train.csv')  
df
```

Out[35]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599 7
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803 5
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450
...
886	887	0	2	Montvila, Rev. Juozas	male	27.0	0	0	211536 1
887	888	1	1	Graham, Miss. Margaret Edith	female	19.0	0	0	112053 3
888	889	0	3	Johnston, Miss. Catherine Helen "Carrie"	female	NaN	1	2	W./C. 6607 2
889	890	1	1	Behr, Mr. Karl Howell	male	26.0	0	0	111369 3
890	891	0	3	Dooley, Mr. Patrick	male	32.0	0	0	370376

891 rows × 12 columns



```
In [45]: median_age = df['Age'].median() # Calculate median of the 'age' column
df['Age'].fillna(median_age, inplace=True) # Replace NaNs with median
df
```

Out[45]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	7
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	5
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	
...	
886	887	0	2	Montvila, Rev. Juozas	male	27.0	0	0	211536	1
887	888	1	1	Graham, Miss. Margaret Edith	female	19.0	0	0	112053	3
888	889	0	3	Johnston, Miss. Catherine Helen "Carrie"	female	28.0	1	2	W./C. 6607	2
889	890	1	1	Behr, Mr. Karl Howell	male	26.0	0	0	111369	3
890	891	0	3	Dooley, Mr. Patrick	male	32.0	0	0	370376	

891 rows × 12 columns



In [46]: `y=df['Survived']`
`y`


```
Out[46]: 0      0
         1      1
         2      1
         3      1
         4      0
         ..
        886    0
        887    1
        888    0
        889    1
        890    0
        Name: Survived, Length: 891, dtype: int64
```

```
In [47]: x = df[['Pclass', 'Sex', 'Age', 'Fare']] # Step 1: Select relevant columns
         x = pd.get_dummies(x, columns=['Sex'], drop_first=True) # Step 2: One-hot encoding
         x['Sex_male'] = x['Sex_male'].astype(int)
         x
```

```
Out[47]:
```

	Pclass	Age	Fare	Sex_male
0	3	22.0	7.2500	1
1	1	38.0	71.2833	0
2	3	26.0	7.9250	0
3	1	35.0	53.1000	0
4	3	35.0	8.0500	1
...
886	2	27.0	13.0000	1
887	1	19.0	30.0000	0
888	3	28.0	23.4500	0
889	1	26.0	30.0000	1
890	3	32.0	7.7500	1

891 rows × 4 columns

```
In [48]: from sklearn.model_selection import train_test_split
         x_train, x_test, y_train, y_test = train_test_split(x,y, test_size=0.30,
                                                             random_state=101)
```

```
In [49]: from sklearn.ensemble import RandomForestClassifier
         from sklearn.metrics import confusion_matrix, accuracy_score, classification_report
```

```
In [50]: rf_model = RandomForestClassifier(random_state=101, n_estimators=100)
```

```
In [51]: rf_model.fit(x_train, y_train)
```

```
Out[51]: ▼      RandomForestClassifier
          RandomForestClassifier(random_state=101)
```

```
In [53]: rf_predictions = rf_model.predict(x_test)
```

```
In [55]: model.score(x_test,y_test)
```

```
Out[55]: 0.7985074626865671
```

```
In [56]: accuracy=confusion_matrix(y_test,predictions)
accuracy
```

```
Out[56]: array([[135, 19],
               [ 35, 79]], dtype=int64)
```

```
In [57]: from sklearn.metrics import classification_report
print(classification_report(y_test,predictions))
```

	precision	recall	f1-score	support
0	0.79	0.88	0.83	154
1	0.81	0.69	0.75	114
accuracy			0.80	268
macro avg	0.80	0.78	0.79	268
weighted avg	0.80	0.80	0.80	268

```
In [ ]: #SVM
```

```
In [58]: df = pd.read_csv('C:/Users/user/Downloads/titanic_train.csv')
df
```

Out[58]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599 7
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803 5
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450
...
886	887	0	2	Montvila, Rev. Juozas	male	27.0	0	0	211536 1
887	888	1	1	Graham, Miss. Margaret Edith	female	19.0	0	0	112053 3
888	889	0	3	Johnston, Miss. Catherine Helen "Carrie"	female	NaN	1	2	W./C. 6607 2
889	890	1	1	Behr, Mr. Karl Howell	male	26.0	0	0	111369 3
890	891	0	3	Dooley, Mr. Patrick	male	32.0	0	0	370376

891 rows × 12 columns



```
In [59]: median_age = df['Age'].median() # Calculate median of the 'age' column
df['Age'].fillna(median_age, inplace=True) # Replace NaNs with median
df
```

Out[59]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	7
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	5
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	
...	
886	887	0	2	Montvila, Rev. Juozas	male	27.0	0	0	211536	1
887	888	1	1	Graham, Miss. Margaret Edith	female	19.0	0	0	112053	3
888	889	0	3	Johnston, Miss. Catherine Helen "Carrie"	female	28.0	1	2	W./C. 6607	2
889	890	1	1	Behr, Mr. Karl Howell	male	26.0	0	0	111369	3
890	891	0	3	Dooley, Mr. Patrick	male	32.0	0	0	370376	

891 rows × 12 columns



In [60]: `y=df['Survived']`
`y`

```
Out[60]: 0      0
         1      1
         2      1
         3      1
         4      0
         ..
        886    0
        887    1
        888    0
        889    1
        890    0
        Name: Survived, Length: 891, dtype: int64
```

```
In [61]: x = df[['Pclass', 'Sex', 'Age', 'Fare']] # Step 1: Select relevant columns
         x = pd.get_dummies(x, columns=['Sex'], drop_first=True) # Step 2: One-hot encoding
         x['Sex_male'] = x['Sex_male'].astype(int)
         x
```

```
Out[61]:
```

	Pclass	Age	Fare	Sex_male
0	3	22.0	7.2500	1
1	1	38.0	71.2833	0
2	3	26.0	7.9250	0
3	1	35.0	53.1000	0
4	3	35.0	8.0500	1
...
886	2	27.0	13.0000	1
887	1	19.0	30.0000	0
888	3	28.0	23.4500	0
889	1	26.0	30.0000	1
890	3	32.0	7.7500	1

891 rows × 4 columns

```
In [62]: from sklearn.model_selection import train_test_split
         x_train, x_test, y_train, y_test = train_test_split(x,y, test_size=0.30,
                                                             random_state=101)
```

```
In [73]: from sklearn.preprocessing import StandardScaler
         scaler = StandardScaler()

         # Fit on training data and transform both train and test
         x_train_scaled = scaler.fit_transform(x_train)
         x_test_scaled = scaler.transform(x_test)
```

```
In [75]: from sklearn.svm import SVC
         from sklearn.metrics import confusion_matrix, accuracy_score, classification_report
```

```
In [76]: svm_model = SVC(kernel='rbf') # You can also try 'rbf', 'poly', or 'sigmoid'
```

```
In [77]: svm_model.fit(x_train_scaled, y_train)
```

```
Out[77]: ▼ SVC
          SVC()
```

```
In [78]: svm_predictions = svm_model.predict(x_test_scaled)
```

```
In [81]: model.score(x_test_scaled, y_test)
```

D:\anaconda\Lib\site-packages\sklearn\base.py:439: UserWarning: X does not have valid feature names, but DecisionTreeClassifier was fitted with feature names
warnings.warn(

```
Out[81]: 0.23134328358208955
```

```
In [80]: accuracy=confusion_matrix(y_test, predictions)
          accuracy
```

```
Out[80]: array([[135, 19],
                [ 35, 79]], dtype=int64)
```

```
In [82]: from sklearn.metrics import classification_report
          print(classification_report(y_test, predictions))
```

	precision	recall	f1-score	support
0	0.79	0.88	0.83	154
1	0.81	0.69	0.75	114
accuracy			0.80	268
macro avg	0.80	0.78	0.79	268
weighted avg	0.80	0.80	0.80	268

```
In [71]: #DECISION TREE
```

```
In [1]: import pandas as pd
          import numpy as np
          import matplotlib.pyplot as plt
          import seaborn as sns
          %matplotlib inline
```

```
In [3]: df = pd.read_csv('C:/Users/user/Downloads/titanic_train.csv')
          df
```

Out[3]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599 7
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803 5
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450
...
886	887	0	2	Montvila, Rev. Juozas	male	27.0	0	0	211536 1
887	888	1	1	Graham, Miss. Margaret Edith	female	19.0	0	0	112053 3
888	889	0	3	Johnston, Miss. Catherine Helen "Carrie"	female	NaN	1	2	W./C. 6607 2
889	890	1	1	Behr, Mr. Karl Howell	male	26.0	0	0	111369 3
890	891	0	3	Dooley, Mr. Patrick	male	32.0	0	0	370376

891 rows × 12 columns



In [5]: `y=df['Survived']`
`y`

```
Out[5]: 0      0
        1      1
        2      1
        3      1
        4      0
        ..
        886    0
        887    1
        888    0
        889    1
        890    0
        Name: Survived, Length: 891, dtype: int64
```

```
In [10]: import pandas as pd

# Assuming your DataFrame is named df
median_age = df['Age'].median() # Calculate median of the 'age' column
df['Age'].fillna(median_age, inplace=True) # Replace NaNs with median
df
```


Out[10]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	7
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	5
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	
...	
886	887	0	2	Montvila, Rev. Juozas	male	27.0	0	0	211536	1
887	888	1	1	Graham, Miss. Margaret Edith	female	19.0	0	0	112053	3
888	889	0	3	Johnston, Miss. Catherine Helen "Carrie"	female	28.0	1	2	W./C. 6607	2
889	890	1	1	Behr, Mr. Karl Howell	male	26.0	0	0	111369	3
890	891	0	3	Dooley, Mr. Patrick	male	32.0	0	0	370376	

891 rows × 12 columns



In [12]: `y=df['Survived']`
`y`

```
Out[12]: 0      0
          1      1
          2      1
          3      1
          4      0
          ..
          886    0
          887    1
          888    0
          889    1
          890    0
          Name: Survived, Length: 891, dtype: int64
```

```
In [21]: x = df[['Pclass', 'Sex', 'Age', 'Fare']] # Step 1: Select relevant columns
          x = pd.get_dummies(x, columns=['Sex'], drop_first=True) # Step 2: One-hot encoding
          x['Sex_male'] = x['Sex_male'].astype(int)
          x
```

```
Out[21]:
```

	Pclass	Age	Fare	Sex_male
0	3	22.0	7.2500	1
1	1	38.0	71.2833	0
2	3	26.0	7.9250	0
3	1	35.0	53.1000	0
4	3	35.0	8.0500	1
...
886	2	27.0	13.0000	1
887	1	19.0	30.0000	0
888	3	28.0	23.4500	0
889	1	26.0	30.0000	1
890	3	32.0	7.7500	1

891 rows × 4 columns

```
In [22]: from sklearn.model_selection import train_test_split
          x_train, x_test, y_train, y_test = train_test_split(x,y, test_size=0.30,
                                                             random_state=101)
```

```
In [23]: from sklearn import tree
          model=tree.DecisionTreeClassifier()
          model.fit(x_train,y_train)
```

```
Out[23]: ▾ DecisionTreeClassifier
          DecisionTreeClassifier()
```

```
In [54]: predictions=model.predict(x_test)
```

```
In [25]: y_test
```

```
Out[25]: 331    0
          700    1
          748    0
          751    1
          481    0
          ..
          388    0
          416    1
          407    1
          482    0
          829    1
          Name: Survived, Length: 268, dtype: int64
```

```
In [26]: model.score(x_test,y_test)
```

```
Out[26]: 0.7985074626865671
```

```
In [27]: from sklearn.metrics import confusion_matrix
```

```
In [31]: accuracy=confusion_matrix(y_test,predictions)
          accuracy
```

```
Out[31]: array([[135,  19],
                [ 35,  79]], dtype=int64)
```

```
In [32]: from sklearn.metrics import classification_report
          print(classification_report(y_test,predictions))
```

	precision	recall	f1-score	support
0	0.79	0.88	0.83	154
1	0.81	0.69	0.75	114
accuracy			0.80	268
macro avg	0.80	0.78	0.79	268
weighted avg	0.80	0.80	0.80	268

```
In [ ]: #NAIVE BAYES
```

```
In [83]: import pandas as pd
          import numpy as np
          import matplotlib.pyplot as plt
          import seaborn as sns
          %matplotlib inline
```

```
In [84]: df = pd.read_csv('C:/Users/user/Downloads/titanic_train.csv')
          df
```

Out[84]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599 7
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803 5
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450
...
886	887	0	2	Montvila, Rev. Juozas	male	27.0	0	0	211536 1
887	888	1	1	Graham, Miss. Margaret Edith	female	19.0	0	0	112053 3
888	889	0	3	Johnston, Miss. Catherine Helen "Carrie"	female	NaN	1	2	W./C. 6607 2
889	890	1	1	Behr, Mr. Karl Howell	male	26.0	0	0	111369 3
890	891	0	3	Dooley, Mr. Patrick	male	32.0	0	0	370376

891 rows × 12 columns



In [85]: `y=df['Survived']`
`y`

```
Out[85]: 0      0
         1      1
         2      1
         3      1
         4      0
         ..
        886     0
        887     1
        888     0
        889     1
        890     0
        Name: Survived, Length: 891, dtype: int64
```

```
In [86]: # Assuming your DataFrame is named df
median_age = df['Age'].median() # Calculate median of the 'age' column
df['Age'].fillna(median_age, inplace=True) # Replace NaNs with median
df
```

Out[86]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	7
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	5
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	
...	
886	887	0	2	Montvila, Rev. Juozas	male	27.0	0	0	211536	1
887	888	1	1	Graham, Miss. Margaret Edith	female	19.0	0	0	112053	3
888	889	0	3	Johnston, Miss. Catherine Helen "Carrie"	female	28.0	1	2	W./C. 6607	2
889	890	1	1	Behr, Mr. Karl Howell	male	26.0	0	0	111369	3
890	891	0	3	Dooley, Mr. Patrick	male	32.0	0	0	370376	

891 rows × 12 columns



```
In [88]: x = df[['Pclass', 'Sex', 'Age', 'Fare']] # Step 1: Select relevant columns
x = pd.get_dummies(x, columns=['Sex'], drop_first=True) # Step 2: One-hot encoding
x['Sex_male'] = x['Sex_male'].astype(int)
x
```

Out[88]:

	Pclass	Age	Fare	Sex_male
0	3	22.0	7.2500	1
1	1	38.0	71.2833	0
2	3	26.0	7.9250	0
3	1	35.0	53.1000	0
4	3	35.0	8.0500	1
...
886	2	27.0	13.0000	1
887	1	19.0	30.0000	0
888	3	28.0	23.4500	0
889	1	26.0	30.0000	1
890	3	32.0	7.7500	1

891 rows × 4 columns

```
In [101... from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x,y, test_size=0.30,
                                                    random_state=101)
```

```
In [103... from sklearn.naive_bayes import GaussianNB

model_nb = GaussianNB()           # ✓ Instantiate the model
model_nb.fit(x_train, y_train)    # ✓ Fit the model
```

```
Out[103... ▼ GaussianNB
GaussianNB()
```

```
In [104... predictions=model.predict(x_test)
```

```
In [105... model.score(x_test,y_test)
```

```
Out[105... 0.7910447761194029
```

```
In [106... from sklearn.metrics import confusion_matrix
accuracy=confusion_matrix(y_test,predictions)
accuracy
```

```
Out[106... array([[135, 19],
       [ 37, 77]], dtype=int64)
```

```
In [107... from sklearn.metrics import classification_report
print(classification_report(y_test,predictions))
```

	precision	recall	f1-score	support
0	0.78	0.88	0.83	154
1	0.80	0.68	0.73	114
accuracy			0.79	268
macro avg	0.79	0.78	0.78	268
weighted avg	0.79	0.79	0.79	268

In [108... `#KNN`

In [123... `import pandas as pd`
`import numpy as np`
`import matplotlib.pyplot as plt`
`import seaborn as sns`
`%matplotlib inline`

In [124... `df = pd.read_csv('C:/Users/user/Downloads/titanic_train.csv')`
`df`

Out[124...

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599 7
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803 5
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450
...
886	887	0	2	Montvila, Rev. Juozas	male	27.0	0	0	211536 1
887	888	1	1	Graham, Miss. Margaret Edith	female	19.0	0	0	112053 3
888	889	0	3	Johnston, Miss. Catherine Helen "Carrie"	female	NaN	1	2	W./C. 6607 2
889	890	1	1	Behr, Mr. Karl Howell	male	26.0	0	0	111369 3
890	891	0	3	Dooley, Mr. Patrick	male	32.0	0	0	370376

891 rows × 12 columns



In [126...

```
y=df['Survived']
y
# Assuming your DataFrame is named df
median_age = df['Age'].median() # Calculate median of the 'age' column
df['Age'].fillna(median_age, inplace=True) # Replace NaNs with median
```

```
df
x = df[['Pclass', 'Sex', 'Age', 'Fare']] # Step 1: Select relevant columns
x = pd.get_dummies(x, columns=['Sex'], drop_first=True) # Step 2: One-hot encoding
x['Sex_male'] = x['Sex_male'].astype(int)
x
```

Out[126...

	Pclass	Age	Fare	Sex_male
0	3	22.0	7.2500	1
1	1	38.0	71.2833	0
2	3	26.0	7.9250	0
3	1	35.0	53.1000	0
4	3	35.0	8.0500	1
...
886	2	27.0	13.0000	1
887	1	19.0	30.0000	0
888	3	28.0	23.4500	0
889	1	26.0	30.0000	1
890	3	32.0	7.7500	1

891 rows × 4 columns

```
In [127... from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x,y, test_size=0.30,
                                                    random_state=101)
```

```
In [128... from sklearn.neighbors import KNeighborsClassifier
model_knn=KNeighborsClassifier(n_neighbors=2)
model_knn.fit(x_train, y_train) # ✓ Fit the model
```

```
Out[128... KNeighborsClassifier
KNeighborsClassifier(n_neighbors=2)
```

```
In [129... predictions=model.predict(x_test)
```

```
In [130... model.score(x_test,y_test)
```

Out[130... 0.832089552238806

```
In [131... from sklearn.metrics import confusion_matrix
accuracy=confusion_matrix(y_test,predictions)
accuracy
```

```
Out[131... array([[143, 11],
       [ 34, 80]], dtype=int64)
```

```
In [132... from sklearn.metrics import classification_report
print(classification_report(y_test,predictions))
```

	precision	recall	f1-score	support
0	0.81	0.93	0.86	154
1	0.88	0.70	0.78	114
accuracy			0.83	268
macro avg	0.84	0.82	0.82	268
weighted avg	0.84	0.83	0.83	268

In [133... `#XG BOOST`

In [134... `df = pd.read_csv('C:/Users/user/Downloads/titanic_train.csv')`
`df`

Out[134...

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599 7
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803 5
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450
...
886	887	0	2	Montvila, Rev. Juozas	male	27.0	0	0	211536 1
887	888	1	1	Graham, Miss. Margaret Edith	female	19.0	0	0	112053 3
888	889	0	3	Johnston, Miss. Catherine Helen "Carrie"	female	NaN	1	2	W./C. 6607 2
889	890	1	1	Behr, Mr. Karl Howell	male	26.0	0	0	111369 3
890	891	0	3	Dooley, Mr. Patrick	male	32.0	0	0	370376

891 rows × 12 columns



In [135...

```

y=df['Survived']
y
# Assuming your DataFrame is named df
median_age = df['Age'].median() # Calculate median of the 'age' column
df['Age'].fillna(median_age, inplace=True) # Replace NaNs with median

```

```
df
x = df[['Pclass', 'Sex', 'Age', 'Fare']] # Step 1: Select relevant columns
x = pd.get_dummies(x, columns=['Sex'], drop_first=True) # Step 2: One-hot encoding
x['Sex_male'] = x['Sex_male'].astype(int)
x
```

Out[135...

	Pclass	Age	Fare	Sex_male
0	3	22.0	7.2500	1
1	1	38.0	71.2833	0
2	3	26.0	7.9250	0
3	1	35.0	53.1000	0
4	3	35.0	8.0500	1
...
886	2	27.0	13.0000	1
887	1	19.0	30.0000	0
888	3	28.0	23.4500	0
889	1	26.0	30.0000	1
890	3	32.0	7.7500	1

891 rows × 4 columns

```
In [136... from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x,y, test_size=0.30,
                                                    random_state=101)
```

```
In [137... !pip install xgboost
from xgboost import XGBClassifier
```

Requirement already satisfied: xgboost in d:\anaconda\lib\site-packages (3.0.1)
Requirement already satisfied: numpy in d:\anaconda\lib\site-packages (from xgboost) (1.26.4)
Requirement already satisfied: scipy in d:\anaconda\lib\site-packages (from xgboost) (1.11.4)

```
In [138... model = XGBClassifier(use_label_encoder=False, eval_metric='logloss', random_state=101)
model.fit(x_train, y_train)

# Predict on test data
y_pred = model.predict(x_test)
```

D:\anaconda\Lib\site-packages\xgboost\training.py:183: UserWarning: [16:32:55] WARNING: C:\actions-runner_work\xgboost\xgboost\src\learner.cc:738: Parameters: { "use_label_encoder" } are not used.

```
bst.update(dtrain, iteration=i, fobj=obj)
```

```
In [139... model.score(x_test,y_test)
```

Out[139... 0.832089552238806

```
In [140]: from sklearn.metrics import confusion_matrix
accuracy=confusion_matrix(y_test,predictions)
accuracy
```

```
Out[140]: array([[143,  11],
               [ 34,  80]], dtype=int64)
```

```
In [141]: from sklearn.metrics import classification_report
print(classification_report(y_test,predictions))
```

	precision	recall	f1-score	support
0	0.81	0.93	0.86	154
1	0.88	0.70	0.78	114
accuracy			0.83	268
macro avg	0.84	0.82	0.82	268
weighted avg	0.84	0.83	0.83	268

```
In [ ]: #ROC AND AUC CURVE FOR XG BOOST(OUR BEST MODEL)
```

```
In [10]: # Step 1: Import required Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, classification_report, roc_curve,
from xgboost import XGBClassifier
```

```
In [11]: # Step 2: Load the dataset
df = pd.read_csv('C:/Users/user/Downloads/titanic_train.csv')
```

```
In [12]: # Step 3: Handle missing values (e.g., Age)
median_age = df['Age'].median()
df['Age'].fillna(median_age, inplace=True)
```

```
In [13]: # Step 4: Define features (X) and target (y)
x = df[['Pclass', 'Sex', 'Age', 'Fare']]
x = pd.get_dummies(x, columns=['Sex'], drop_first=True) # Convert 'Sex' to nume
y = df['Survived']
```

```
In [14]: # Step 5: Split the data
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.3, random_
```

```
In [15]: # Step 6: Train the XGBoost model
model = XGBClassifier(use_label_encoder=False, eval_metric='logloss', random_sta
model.fit(x_train, y_train)
```

D:\anaconda\Lib\site-packages\xgboost\training.py:183: UserWarning: [20:05:05] WA
ARNING: C:\actions-runner_work\xgboost\xgboost\src\learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

```
bst.update(dtrain, iteration=i, fobj=obj)
```

Out[15]:

```

XGBClassifier
XGBClassifier(base_score=None, booster=None, callbacks=None,
               colsample_bylevel=None, colsample_bynode=None,
               colsample_bytree=None, device=None, early_stopping_rounds=None,
               enable_categorical=False, eval_metric='logloss',
               feature_types=None, feature_weights=None, gamma=None,
               grow_policy=None, importance_type=None,
               interaction_constraints=None, learning_rate=None, max_bin=None,

```

```

In [16]: # Step 7: Predictions and evaluation
y_pred = model.predict(x_test)

# Confusion matrix
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))

# Classification report
print("\nClassification Report:")
print(classification_report(y_test, y_pred))

```

Confusion Matrix:

```
[[143  11]
 [ 34  80]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.81	0.93	0.86	154
1	0.88	0.70	0.78	114
accuracy			0.83	268
macro avg	0.84	0.82	0.82	268
weighted avg	0.84	0.83	0.83	268

```

In [17]: # Get predicted probabilities for the positive class (Survived=1)
y_prob = model.predict_proba(x_test)[:, 1]

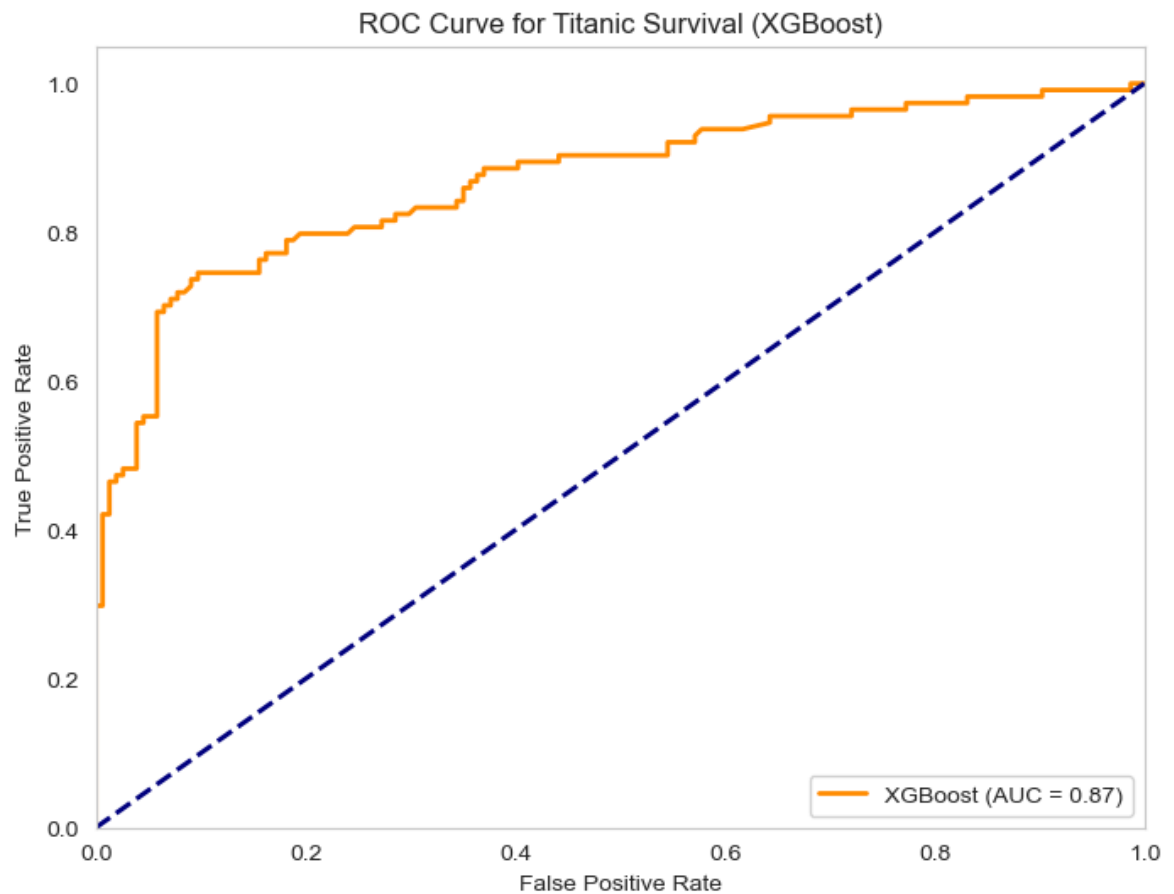
# Calculate FPR, TPR, thresholds
fpr, tpr, thresholds = roc_curve(y_test, y_prob)

# Compute AUC
roc_auc = auc(fpr, tpr)

# Plot ROC Curve
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='darkorange', lw=2, label='XGBoost (AUC = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--') # Diagonal line
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve for Titanic Survival (XGBoost)')
plt.legend(loc="lower right")

```

```
plt.grid()  
plt.show()
```



```
In [18]: print("AUC Score:", roc_auc)
```

AUC Score: 0.8703292321713374